

DSC680T301_PhishingDetection_LincolnBrown

February 28, 2025

```
[52]: from collections import Counter
import matplotlib.pyplot as plt
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
import numpy as np
import pandas as pd
import re
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix, roc_auc_score, roc_curve
import torch
from torch.utils.data import DataLoader, TensorDataset
from torch.optim import AdamW
import torch.nn.functional as F
from transformers import BertTokenizer, BertForSequenceClassification, \
    get_scheduler
from textblob import TextBlob
```

```
[2]: f = './Phishing_Email.csv'
df = pd.read_csv(f, index_col=0)
```

```
[3]: df
```

```
[3]:
```

	Email Text	Email Type
0	re : 6 . 1100 , disc : uniformitarianism , re ...	Safe Email
1	the other side of * galicismos * * galicismo *...	Safe Email
2	re : equistar deal tickets are you still avail...	Safe Email
3	\nHello I am your hot lil horny toy.\n I am...	Phishing Email
4	software at incredibly low prices (86 % lower...	Phishing Email
...
18646	date a lonely housewife always wanted to date ...	Phishing Email
18647	request submitted : access request for anita ...	Safe Email
18648	re : important - prc mtg hi dorn & john , as y...	Safe Email
18649	press clippings - letter on californian utilit...	Safe Email

18650

empty Phishing Email

[18650 rows x 2 columns]

0.1 Data Cleaning

```
[4]: # Get the number of null values in the dataset
df.isnull().sum()
```

```
[4]: Email Text      16
     Email Type      0
     dtype: int64
```

```
[5]: # Drop the na's
clean_df = df.copy()
clean_df.dropna(inplace=True)
```

```
[6]: # Look for records that contain the text 'empty'
     # These records will also be considered missing and dropped as well

len(clean_df.loc[clean_df['Email Text'] == 'empty'])
```

```
[6]: 533
```

```
[7]: # Select records that do not have the email text 'empty'
clean_df = clean_df[clean_df['Email Text'] != 'empty']
```

```
[8]: # View the shape after removing nulls
clean_df.shape
```

```
[8]: (18101, 2)
```

```
[9]: # Binary Encode the Email Type
clean_df.loc[:, 'Email Type'] = clean_df['Email Type'].map({'Phishing Email': 1,
    ↪1, 'Safe Email': 0})
```

```
[10]: # Convert to int64
clean_df['Email Type'] = pd.to_numeric(clean_df['Email Type'])
clean_df['Email Type'].dtype
```

```
[10]: dtype('int64')
```

```
[11]: # Download stopwords from nltk
nltk.download('stopwords')
nltk.download('punkt')
```

[nltk_data] Downloading package stopwords to /Users/x/nltk_data...

[nltk_data] Package stopwords is already up-to-date!

```
[nltk_data] Downloading package punkt to /Users/x/nltk_data...  
[nltk_data]   Package punkt is already up-to-date!
```

```
[11]: True
```

0.2 Exploratory Data Analysis

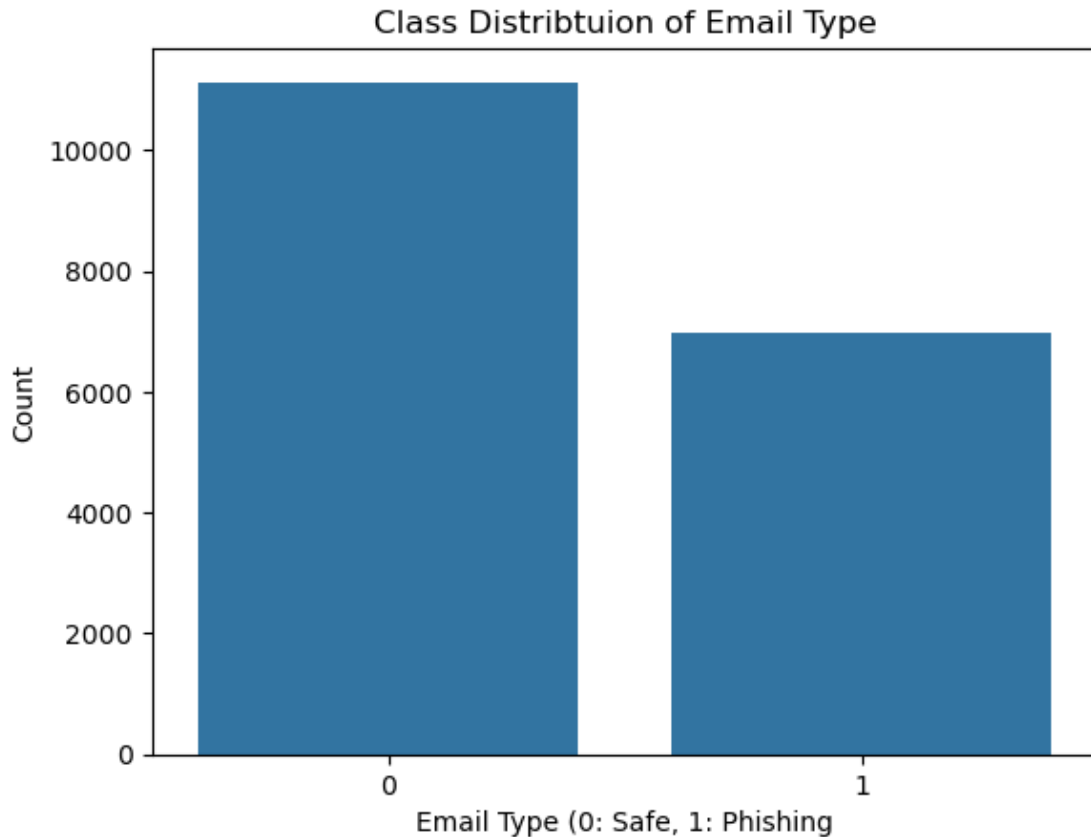
```
[12]: eda_df = clean_df.copy()
```

```
[13]: eda_df.describe()
```

```
[13]:
```

	Email Type
count	18101.000000
mean	0.385448
std	0.486714
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	1.000000

```
[14]: # Create a count plot of the class distribution  
sns.countplot(x=eda_df['Email Type'])  
plt.title('Class Distribtuion of Email Type')  
plt.xlabel('Email Type (0: Safe, 1: Phishing)')  
plt.ylabel('Count')  
plt.show()
```



```
[15]: # Get the word and character counts for the emails

eda_df['word_count'] = eda_df['Email Text'].apply(lambda x: len(x.split()))
eda_df['char_count'] = eda_df['Email Text'].apply(lambda x: len(x))
```

```
[16]: eda_df.describe()
```

```
[16]:
```

	Email Type	word_count	char_count
count	18101.000000	1.810100e+04	1.810100e+04
mean	0.385448	5.509090e+02	2.836649e+03
std	0.486714	2.622834e+04	1.266919e+05
min	0.000000	0.000000e+00	1.000000e+00
25%	0.000000	8.000000e+01	4.340000e+02
50%	0.000000	1.670000e+02	9.130000e+02
75%	1.000000	3.660000e+02	1.936000e+03
max	1.000000	3.527576e+06	1.703669e+07

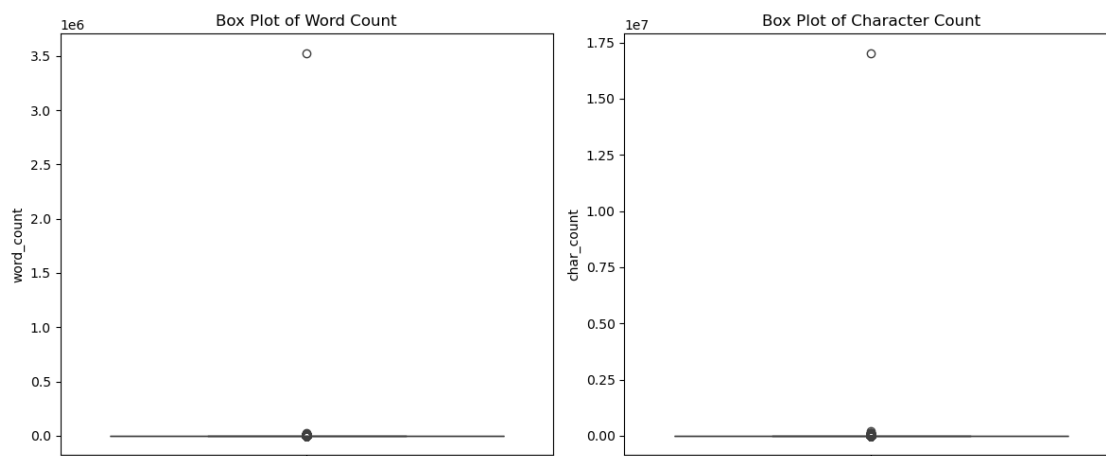
From the descriptive statistics above, we can see that we have at least one email that contains no characters and at least one that is many magnitudes higher ($e+07$) than the median ($e+02$). Let's take a look at the distribution of `word_count` and `char_count`.

```
[17]: # Review the distributions using box plots
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

# Box plot of word_count
sns.boxplot(eda_df['word_count'], ax=axes[0])
axes[0].set_title('Box Plot of Word Count')

# Box plot of char_count
sns.boxplot(eda_df['char_count'], ax=axes[1])
axes[1].set_title('Box Plot of Character Count')

plt.tight_layout()
plt.show()
```



```
[18]: # Remove the outliers using IQR
q1 = eda_df['char_count'].quantile(0.25)
q3 = eda_df['char_count'].quantile(0.75)
iqr = q3 - q1

# Define bounds
lower_bound = q1 - 1.5 * iqr
upper_bound = q3 + 1.5 * iqr

# Filter the outliers
filtered_df = eda_df[(eda_df['char_count'] >= lower_bound) &
                    (eda_df['char_count'] <= upper_bound)].copy()
```

```
[19]: # Review the distributions after removing outliers
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

# Box plot of word_count
```

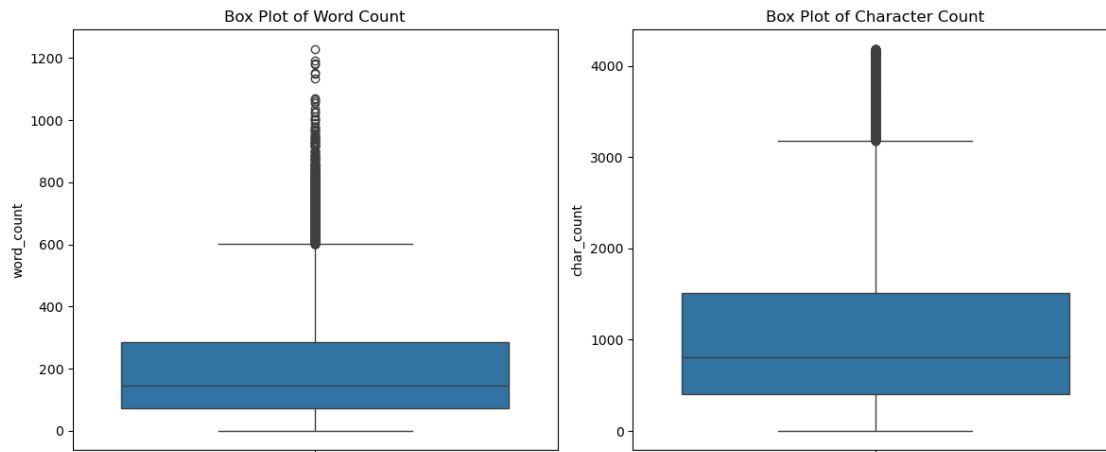
```

sns.boxplot(filtered_df['word_count'], ax=axes[0])
axes[0].set_title('Box Plot of Word Count')

# Box plot of char_count
sns.boxplot(filtered_df['char_count'], ax=axes[1])
axes[1].set_title('Box Plot of Character Count')

plt.tight_layout()
plt.show()

```



```

[20]: # Determine the amount of outlier records removed
num_outliers_removed = eda_df.shape[0] - filtered_df.shape[0]

print(f'There were {num_outliers_removed} outlier records removed.')
print(f'The dataset now contains {filtered_df.shape[0]} rows and {filtered_df.
↳shape[1]} columns.')

```

There were 1662 outlier records removed.
The dataset now contains 16439 rows and 4 columns.

```

[21]: # Check distributions of safe vs phishing emails for word and character counts

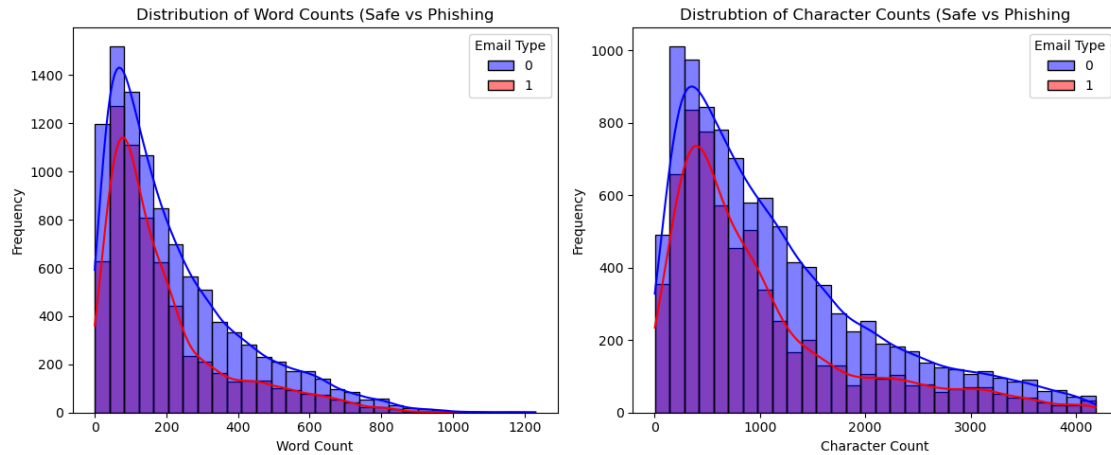
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

# Plot word count
sns.histplot(data=filtered_df, x='word_count', hue='Email Type', bins=30,
↳kde=True, ax=axes[0], palette={0: 'blue', 1: 'red'})
axes[0].set_title('Distribution of Word Counts (Safe vs Phishing)')
axes[0].set_xlabel('Word Count')
axes[0].set_ylabel('Frequency')

```

```
# Plot character count
sns.histplot(data=filtered_df, x='char_count', hue='Email Type', bins=30,
             kde=True, ax=axes[1], palette={0: 'blue', 1: 'red'})
axes[1].set_title('Distrubtion of Character Counts (Safe vs Phishing)')
axes[1].set_xlabel('Character Count')
axes[1].set_ylabel('Frequency')

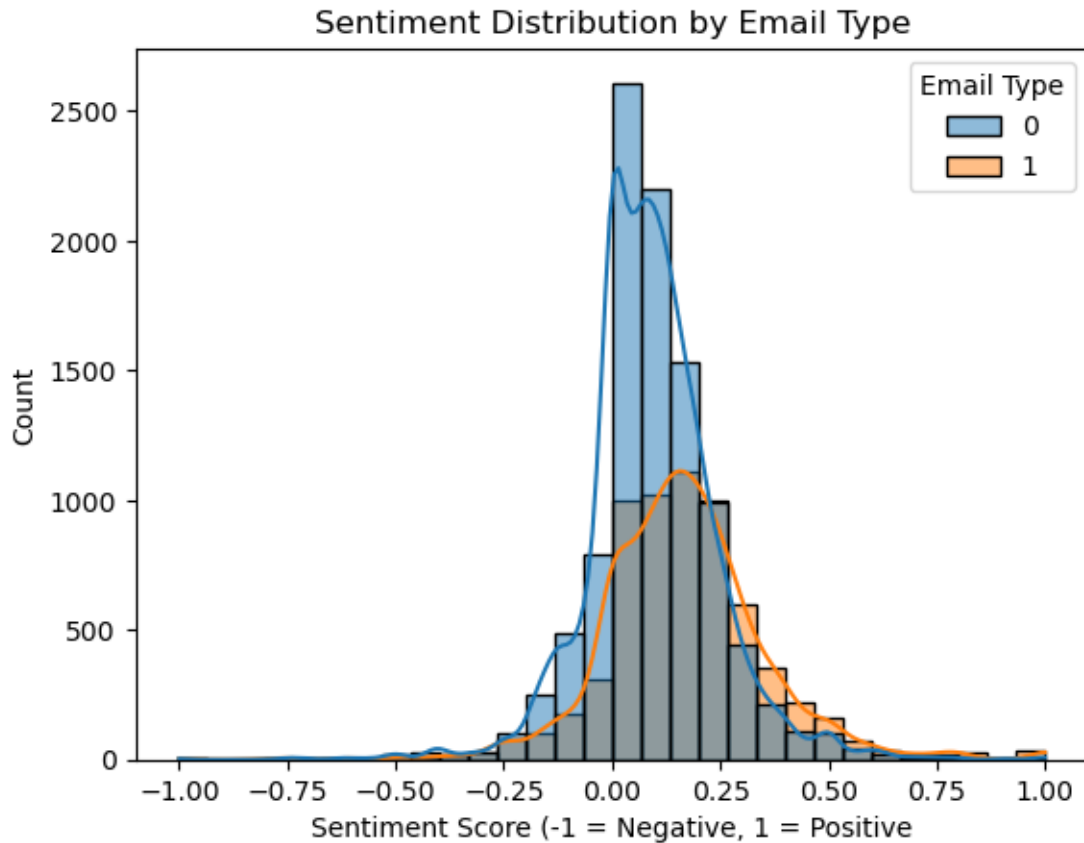
plt.tight_layout()
plt.show()
```



0.3 Sentiment Analysis

```
[22]: filtered_df['sentiment'] = filtered_df.loc[:, 'Email Text'].apply(lambda x:
    ↪ TextBlob(x).sentiment.polarity)

[23]: # Plot sentiment distribution
sns.histplot(filtered_df, x='sentiment', hue='Email Type', bins=30, kde=True)
plt.title('Sentiment Distribution by Email Type')
plt.xlabel('Sentiment Score (-1 = Negative, 1 = Positive)')
plt.show()
```



0.4 Most Common Words in Safe vs Phishing Emails

```
[24]: # Split phishing and safe emails
safe_text = ' '.join(filtered_df[filtered_df['Email Type'] == 0]['Email Text'])
phishing_text = ' '.join(filtered_df[filtered_df['Email Type'] == 1]['Email_
↳Text'])
```

0.5 Bag of Words

```
[25]: # Setup the function to clean the text

def bow_preprocess_text(text):
    text = re.sub(r'\W', ' ', text) # Remove special characters
    text = re.sub(r'\d+', ' ', text) # Remove numbers
    text = re.sub(r'\s_', ' ', text).strip() # Remove extra spaces
    text = text.lower()
    words = word_tokenize(text) # Tokenization
    words = [word for word in words if word not in stopwords.words('english')]
    ↳ # Remove stop words
```



```
return ' '.join(words)
```

```
[26]: model_df = filtered_df.copy()
model_df
```

```
[26]:
```

	Email Text	Email Type	\
0	re : 6 . 1100 , disc : uniformitarianism , re ...	0	
1	the other side of * galicismos * * galicismo *...	0	
2	re : equistar deal tickets are you still avail...	0	
3	\nHello I am your hot lil horny toy.\n I am...	1	
4	software at incredibly low prices (86 % lower...	1	
...	
18643	URL: http://diveintomark.org/archives/2002/09/...	0	
18646	date a lonely housewife always wanted to date ...	1	
18647	request submitted : access request for anita ...	0	
18648	re : important - prc mtg hi dorn & john , as y...	0	
18649	press clippings - letter on californian utilit...	0	

	word_count	char_count	sentiment
0	230	1030	0.201493
1	91	479	0.009375
2	305	1245	0.091540
3	96	688	0.260069
4	91	441	0.152579
...
18643	62	429	0.000000
18646	45	237	-0.100000
18647	99	477	-0.183333
18648	253	1214	0.141250
18649	34	213	0.200000

[16439 rows x 5 columns]

```
[27]: # Preprocess text
model_df['Cleaned Text'] = model_df['Email Text'].apply(bow_preprocess_text)
```

```
[28]: model_df.drop(columns=['Email Text'], inplace=True)
```

```
[29]: model_df
```

```
[29]:
```

	Email Type	word_count	char_count	sentiment	\
0	0	230	1030	0.201493	
1	0	91	479	0.009375	
2	0	305	1245	0.091540	
3	1	96	688	0.260069	
4	1	91	441	0.152579	
...	

18643	0	62	429	0.000000
18646	1	45	237	-0.100000
18647	0	99	477	-0.183333
18648	0	253	1214	0.141250
18649	0	34	213	0.200000

```

                                Cleaned Text
0      disc uniformitarianism sex lang dick hudson ob...
1      side galicismos galicismo spanish term names i...
2      equistar deal tickets still available assist r...
3      hello hot lil horny toy one dream open minded ...
4      software incredibly low prices lower drapery s...
...
18643  url http diveintomark org archives html plan d...
18646  date lonely housewife always wanted date lonel...
18647  request submitted access request anita dupont ...
18648  important prc mtg hi dorn john discovered rece...
18649  press clippings letter californian utilities p...

```

[16439 rows x 5 columns]

```

[30]: # Split the dataset
X_train, X_test, y_train, y_test = train_test_split(model_df['Cleaned Text'],
    ↪model_df['Email Type'], test_size=0.2, random_state=42,
    ↪stratify=model_df['Email Type'])

```

```

[31]: # Convert to TF-IDF features
vectorizer = TfidfVectorizer(max_features=5000)
X_train_tfidf = vectorizer.fit_transform(X_train)
X_test_tfidf = vectorizer.transform(X_test)

```

0.6 Logistic Regression Model

```

[32]: # Logistic Regression Model
# Create the LR model
lr_model = LogisticRegression()
# Train the LR Model
lr_model.fit(X_train_tfidf, y_train)

# Make predictions
y_pred_lr = lr_model.predict(X_test_tfidf)

# Evalute model performance
accuracy = accuracy_score(y_test, y_pred_lr)
print(f'Accuracy: {accuracy:.4f}')
print(f'Classification Report: \n{classification_report(y_test, y_pred_lr)}')

```

Accuracy: 0.9775

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.99	0.98	2008
1	0.98	0.96	0.97	1280
accuracy			0.98	3288
macro avg	0.98	0.97	0.98	3288
weighted avg	0.98	0.98	0.98	3288

```
[33]: feature_names = vectorizer.get_feature_names_out()
top_coefs = lr_model.coef_[0].argsort()[::-1][:20]
print([feature_names[i] for i in top_coefs])


['remove', 'click', 'sightings', 'free', 'money', 'site', 'email', 'save',
'life', 'removed', 'reply', 'hello', 'software', 'offer', 'want', 'viagra',
'best', 'online', 'quality', 'account']

[34]: # Confusion Matrix
lr_cm = confusion_matrix(y_test, y_pred_lr)
lr_cm

[34]: array([[1982,  26],
        [ 48, 1232]])

[65]: # Plot the LR model confusion matrix
plt.figure(figsize=(7,5))
annotations = [f'TN: {lr_cm[0,0]}', f'FP: {lr_cm[0,1]}'], [f'FN:{lr_cm[1,0]}',
↳ f'TP: {lr_cm[1,1]}']
lr_hm = sns.heatmap(lr_cm, annot=True, fmt='d', cmap='Reds',
↳ xticklabels=['Non-Fraud', 'Fraud'], yticklabels=['Non-Fraud', 'Fraud'])

plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Logistic Regression Confusion Matrix')
bert_hm.text(0.1, 0.2, 'True Negative', color='white')
bert_hm.text(1.1, 0.2, 'False Positive')
bert_hm.text(0.1, 1.2, 'False Negative')
bert_hm.text(1.1, 1.2, 'True Positive')
plt.show()
```



```
[36]: # Get the ROC AUC
      # Get the probability for the phishing emails
      y_probs = lr_model.predict_proba(X_test_tfidf)[: , 1]

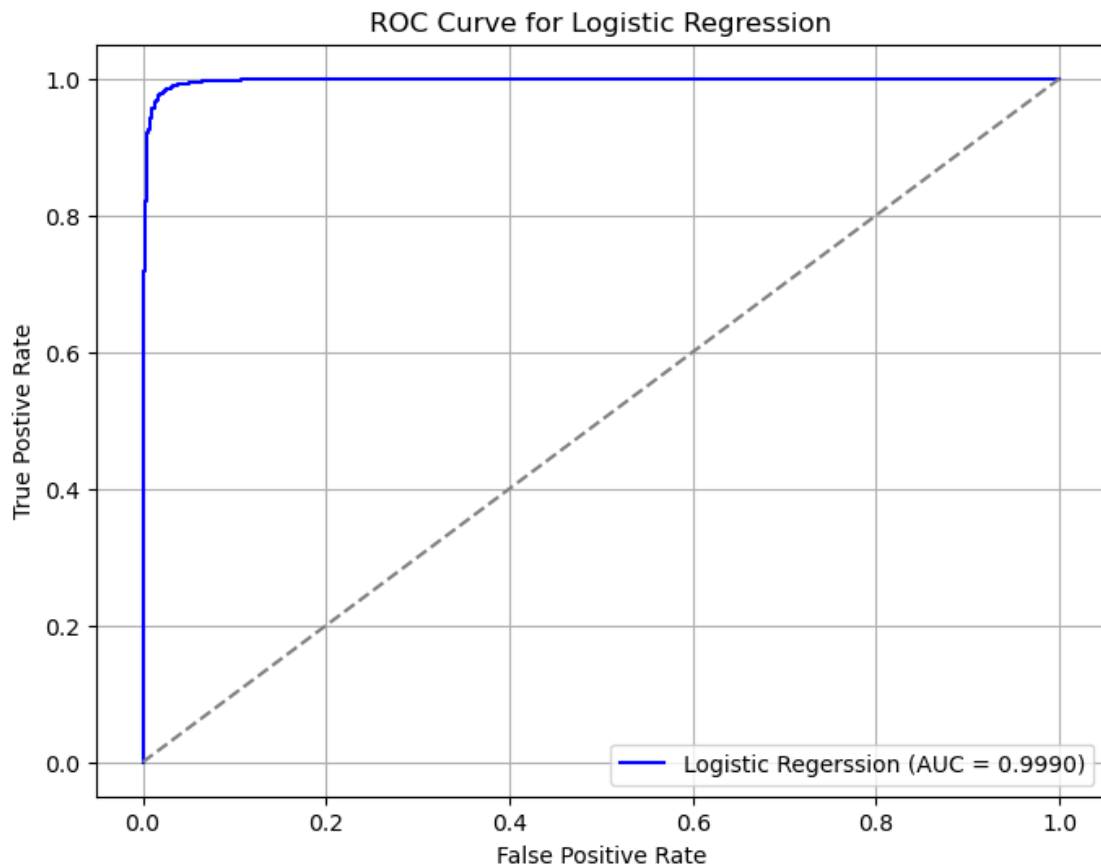
      roc_auc = roc_auc_score(y_test, y_probs)
      roc_auc
```

```
[36]: 0.9973800112051793
```

```
[54]: # Plot the ROC curve
      lr_fpr, lr_tpr, _ = roc_curve(y_test, y_probs)

      plt.figure(figsize=(8,6))
      plt.plot(lr_fpr, lr_tpr, label=f'Logistic Regerssion (AUC = {roc_auc:.4f})',
               color='blue')
      plt.plot([0,1], [0,1], linestyle='--', color='gray')
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Postive Rate')
      plt.title('ROC Curve for Logistic Regression')
```

```
plt.legend()
plt.grid()
plt.show()
```



0.7 BERT Model

```
[38]: # Load the BERT Tokenizer
tokenizer = BertTokenizer.from_pretrained("bert-base-uncased")
```

```
[39]: # Define tokenize function
def tokenize_function(texts):
    return tokenizer(texts.tolist(), padding=True, truncation=True,
        ↪max_length=512, return_tensors='pt')
```

```
[40]: # Tokenize teh data and convert labels to tensors
train_encodings = tokenize_function(X_train)
test_encodings = tokenize_function(X_test)

train_labels = torch.tensor(y_train.tolist())
```

```
test_labels = torch.tensor(y_test.tolist())
```

```
[41]: # Create the Bert Model
bert_model = BertForSequenceClassification.from_pretrained('bert-base-uncased',
↳ num_labels=2)
```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized:

```
['classifier.bias', 'classifier.weight']
```

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
[48]: # Create a TensorDataset for the input_ids and attention masks
train_dataset = TensorDataset(train_encodings['input_ids'],
↳ train_encodings['attention_mask'], train_labels)
test_dataset = TensorDataset(test_encodings['input_ids'],
↳ test_encodings['attention_mask'], test_labels)

# Create data loaders for batch processing
train_dataloader = DataLoader(train_dataset, batch_size=8, shuffle=True)
test_dataloader = DataLoader(test_dataset, batch_size=8, shuffle=False)
```

```
[43]: # Use Apple GPU
device = torch.device('mps' if torch.backends.mps.is_available() else 'cpu')
bert_model.to(device)
```

```
[43]: BertForSequenceClassification(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(30522, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (token_type_embeddings): Embedding(2, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (encoder): BertEncoder(
      (layer): ModuleList(
        (0-11): 12 x BertLayer(
          (attention): BertAttention(
            (self): BertSdpaSelfAttention(
              (query): Linear(in_features=768, out_features=768, bias=True)
              (key): Linear(in_features=768, out_features=768, bias=True)
              (value): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
            (output): BertSelfOutput(
              (dense): Linear(in_features=768, out_features=768, bias=True)
              (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
```

```

        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(intermediate): BertIntermediate(
  (dense): Linear(in_features=768, out_features=3072, bias=True)
  (intermediate_act_fn): GELUActivation()
)
(output): BertOutput(
  (dense): Linear(in_features=3072, out_features=768, bias=True)
  (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
  (dropout): Dropout(p=0.1, inplace=False)
)
)
)
(pooler): BertPooler(
  (dense): Linear(in_features=768, out_features=768, bias=True)
  (activation): Tanh()
)
)
(dropout): Dropout(p=0.1, inplace=False)
(classifier): Linear(in_features=768, out_features=2, bias=True)
)

```

```

[47]: # Train the BERT model
# Set optimizer
optimizer = AdamW(bert_model.parameters(), lr=3e-5)

epochs = 3

num_training_steps = epochs * len(train_dataloader)
lr_scheduler = get_scheduler(
    "linear", optimizer=optimizer, num_warmup_steps=0,
    ↪ num_training_steps=num_training_steps
)

for epoch in range(epochs):
    bert_model.train()
    total_loss = 0
    for batch in train_dataloader:
        input_ids, attention_mask, labels = [x.to(device) for x in batch]
        optimizer.zero_grad()
        outputs = bert_model(input_ids, attention_mask=attention_mask,
        ↪ labels=labels)
        loss = outputs.loss
        total_loss += loss.item()
        loss.backward()

```

```
optimizer.step()

print(f'Epoch {epoch+1}, Loss: {total_loss:.4f}')
```

Epoch 1, Loss: 229.5264

Epoch 2, Loss: 64.3780

Epoch 3, Loss: 31.5956

```
[49]: # Get the BERT model results
bert_model.eval()
predictions, true_labels = [], []

with torch.no_grad():
    for batch in test_dataloader:
        input_ids, attention_mask, labels = [x.to(device) for x in batch]

        outputs = bert_model(input_ids, attention_mask=attention_mask)
        logits = outputs.logits
        preds = torch.argmax(logits, dim=1).cpu().numpy()

        predictions.extend(preds)
        true_labels.extend(labels.cpu().numpy())

# Print metrics
print(f'Accuracy: {accuracy_score(true_labels, predictions):.4f}')
print(classification_report(true_labels, predictions))
```

Accuracy: 0.9857

	precision	recall	f1-score	support
0	0.99	0.99	0.99	2008
1	0.98	0.98	0.98	1280
accuracy			0.99	3288
macro avg	0.99	0.98	0.98	3288
weighted avg	0.99	0.99	0.99	3288

```
[50]: # BERT Model Confusion Matrix
bert_cm = confusion_matrix(true_labels, predictions)


bert_cm
```

```
[50]: array([[1986,  22],
           [ 25, 1255]])
```



```
[63]: # Plot the BERT model confusion matrix
plt.figure(figsize=(7,5))
annotations = [f'TN: {bert_cm[0,0]}', f'FP: {bert_cm[0,1]}'], [f'FN:
    ↳{bert_cm[1,0]}', f'TP: {bert_cm[1,1]}']
bert_hm = sns.heatmap(bert_cm, annot=True, fmt='d', cmap='Reds',
    ↳xticklabels=['Non-Fraud', 'Fraud'], yticklabels=['Non-Fraud', 'Fraud'])

plt.xlabel('Predicted')
plt.ylabel('Acutal')
plt.title('BERT Model Confusion Matrix')
bert_hm.text(0.1, 0.2, 'True Negative', color='white')
bert_hm.text(1.1, 0.2, 'False Positive')
bert_hm.text(0.1, 1.2, 'False Negative')
bert_hm.text(1.1, 1.2, 'True Positive')
plt.show()
```



```
[53]: # Convert logits to probabilities using softmax

with torch.no_grad():
```

```

all_logits = []
all_true_labels = []
for batch in test_dataloader:
    input_ids, attention_mask, labels = [x.to(device) for x in batch]
    outputs = bert_model(input_ids, attention_mask=attention_mask)
    logits = outputs.logits
    probs = F.softmax(logits, dim=1)[: , 1].cpu().numpy() # Probabilities
    ↪for the positive class
    all_logits.extend(probs)
    all_true_labels.extend(labels.cpu().numpy())

# Compute AUC-ROC Score
roc_auc = roc_auc_score(all_true_labels, all_logits)
print(f'ROC AUC Score: {roc_auc:.4f}')

```

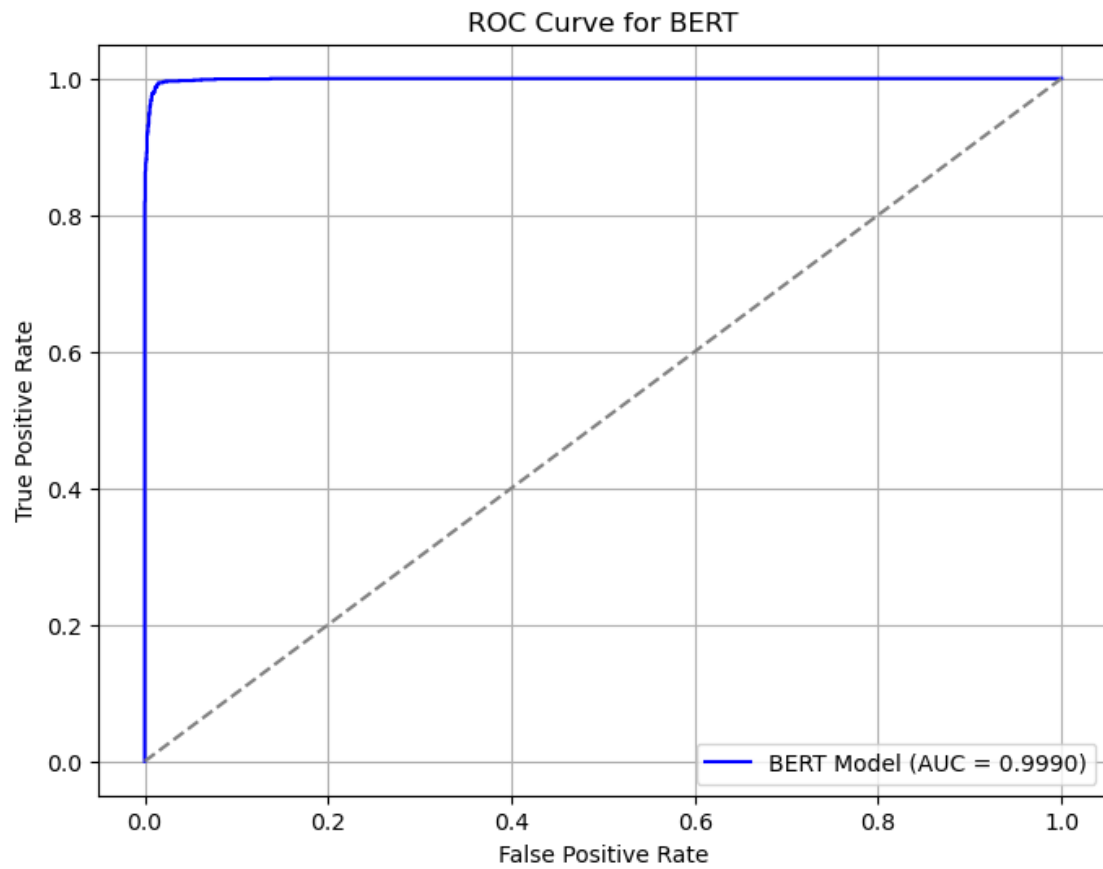
ROC AUC Score: 0.9990

```

[57]: # Plot the ROC curve
bert_fpr, bert_tpr, _ = roc_curve(all_true_labels, all_logits)

plt.figure(figsize=(8,6))
plt.plot(bert_fpr, bert_tpr, label=f'BERT Model (AUC = {roc_auc:.4f})',
    ↪color='blue')
plt.plot([0,1], [0,1], linestyle='--', color='gray')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for BERT')
plt.legend()
plt.grid()
plt.show()

```



[]: